Fantasy Premier League Points Forecasting and Squad Optimization

Project status report

Generated 2026-07-04

1 Conclusion and status

The system - a per-position machine-learning forecaster feeding a MILP squad optimizer - beats its thesis-era LSTM predecessor by roughly 22% in realized backtest points over an identical 31-gameweek window (1,870 vs 1,526 under honest evaluation discipline; Section 5.1). The current state of play, most decision-relevant first:

1. **Hyperparameter tuning delivered the largest single gain in the project's history**: 30 Optuna trials per position (time-ordered CV, tuned strictly on pre-backtest data) lifted the honest backtest from 1,856 to **2,107 realized points** (+13.5%), past even the old leakage-inflated headline. The hand-set defaults were learning too fast and too deep (Section 5.1).
2. **The honest re-baseline first lowered the headline.** Re-running the backtest with combination weights fit strictly before the window showed the long-quoted 1,966 was flattered by weight leakage: the clean untuned figures are 1,869 (NNLS blend) and 1,856 (CatBoost-only) - a statistical tie (Section 5.1). With tuning, **2,107 is the number every future change must beat**.
3. **CatBoost (MAE loss) is the production forecaster at every position**, selected empirically: a 20-gameweek walk-forward head-to-head and a static-window combination bake-off both found it beats the NNLS blend, equal-weight top-k, and ridge stacking on forecast accuracy everywhere (weighted walk-forward MASE 0.684 vs 0.761) - the classic Clemen (1989) result. But that 10% accuracy edge bought *zero* realized points (the tie above), so the honest claim is "equal squads for one-twelfth the training cost", not "better squads" (Section 5.3).
4. **Forecast accuracy is provably not the decision metric.** Twice now, a MASE improvement failed to deliver realized points (fixture/minutes features: 1,966 → 1,880; CatBoost vs blend: tie), and a theory-driven fix - recalibrating CatBoost's median-flattened level before the MILP - *cost* 56 points when tested, because the per-position scalars reshuffle cross-position budget allocation (Section 5.3). Realized-points backtests, plus the new ranking/calibration diagnostics, are the arbiter now; MASE alone selects the wrong models.
5. **No simple or econometric technique tested beats the ML models** (nine baselines, Section 5.2); a further 37 engineered features (opponent form, exponentially-weighted form, per-90 rates, xP) were roughly accuracy-neutral - reported honestly, kept pending hyperparameter tuning, which remains unexplored territory.
6. A code-and-concept review found and fixed four methodological errors (blend-weight leakage, two live-mode staleness bugs, an evaluation asymmetry); the re-baseline above closes out the last of its follow-ups (Section 6).

Priorities, in order: retrain production on the tuned parameters and tune the other GBMs so the registry ranking is default-luck-free (a tuned blend deserves a re-match); stability-check the 2,107 (more trials, second window); probabilistic-upside captaincy; then the larger architectural bets (component decomposition of the target, a global sequence model).

2 Introduction

The system predicts FPL player points one to several gameweeks ahead and converts forecasts into weekly squad decisions (transfers, lineup, captaincy, chips) via a MILP following Kristiansen et al. (2018). The governing principle is honest empirical validation: every candidate technique is tested against the production system on a fixed backtest window, and negative results are reported rather than discarded.

3 Data

Per-gameweek player statistics come from Vaastav Anand's public mirror of the official FPL API (`vaastav/Fantasy-Premier-League`), spanning six seasons - 2020-21 through 2025-26, 162,981 player-gameweek rows. The two oldest seasons lack the Opta expected-goals family and `starts` (introduced 2022-23); these are treated as missing, which the tree-based models handle natively. Gameweeks carry a single ascending index (`GW_global`): dataset season N occupies gameweeks $(N-1) \times 38 + 1$ through $N \times 38$, so the fixed validation window - 2024-25 season, gameweeks 1-31 - is GW153-183.

Exploratory analysis (`notebooks/eda.ipynb`, all six seasons) establishes the two facts the methodology is built on: `total_points` is heavily right-skewed and zero-inflated at every position (Shapiro-Wilk rejects normality, $p < 0.001$ throughout), motivating a scale-free error metric and models free of Gaussian assumptions; and the distribution is stable across seasons, with one caveat - the average defender score-per-gameweek series fails an Augmented Dickey-Fuller stationarity test over the full six seasons ($p = 0.41$; the other positions pass), plausibly reflecting the 2025-26 `defensive_contribution` scoring change.

4 Methodology

4.1 Feature engineering

Each player-gameweek row carries 115 engineered features:

- **Form.** Rolling 3- and 5-gameweek means, previous-gameweek values, and two expanding-mean horizons - season-to-date (resets each season) and career-to-date - over 18 per-gameweek statistics (points, minutes, xG family, BPS, ICT, price, ownership), plus exponentially-weighted means (halflife 3 gameweeks) for the ten most form-sensitive statistics, motivated by the earlier finding that exponential smoothing beats flat windows among simple baselines. All shifted one gameweek relative to the target row, so no row's features contain its own outcome.
- **Rates.** Per-90-minute rates (goals, assists, xG, xA over rolling minutes) separating production efficiency from playing time.
- **Opponent strength.** The upcoming opponent's rolling 6-gameweek attacking/defensive form and clean-sheet rate, computed at team level and merged by opponent with the same strict one-gameweek shift - a dynamic, in-season complement to the static FDR.
- **Fixture difficulty.** The official FPL Fixture Difficulty Rating (1-5) of that gameweek's opponent, plus the mean over this and the next two fixtures. Deliberately *not* shifted: fixture lists are published ahead of play, so these are known-ahead inputs, not leakage.

- **Minutes projection.** Rolling 5-game start rate and 60+-minute rate, shifted like form - a "will he actually play?" signal.
- **FPL's own forecast.** Shifted forms of xP, FPL's published pre-match expected points.

Statistics absent for entire early seasons (the Opta xG family and starts before 2022-23) are kept as missing values rather than zero-filled: "not collected" is a different fact from "recorded zero", and the gradient-boosted models branch on missingness natively.

4.2 Forecasting models

Four independent models, one per position (GK/DEF/MID/FWD) - the statistics that predict a goalkeeper's points are largely disjoint from a forward's, and pooling risks one position's scale dominating the loss. Per position, every regressor in a fixed registry is trained on the same features: LightGBM, XGBoost, CatBoost (MAE loss), OLS, Ridge, ElasticNet, PLS, Random Forest, Extra Trees, k-NN, LinearSVR, and a sample-capped RBF SVR. Boosted models take raw features including missing values; linear/kernel/distance models get a zero-imputer and standardizer. Registry members that tests find unhelpful are retained - the blend simply assigns them near-zero weight, preserving the comparison for later revisiting.

How the registry members are *combined* is itself an empirical question, answered per position by a combination bake-off scored on identical held-out rows: the best single member (no combination at all) vs. non-negative least squares (NNLS) vs. an equal-weight average of the top-k members vs. ridge-regularised stacking - the latter two being the forecast-combination literature's standard remedies for the instability of estimated weights (Clemen, 1989). Whatever wins becomes the production combiner, so the choice tracks the evidence rather than an assumption. Weight fitting is separated by purpose: *evaluation* weights are fit on one half of the held-out test window and scored on the other half, so reported accuracy is a genuine holdout figure; *production and backtest* weights are fit on a 16-gameweek window strictly before whatever is subsequently predicted, so no weight ever sees a gameweek it will be used to forecast. Plain OLS is designated the *index* - the simple benchmark every technique must beat, as a passive market index is the bar an active strategy must clear.

Beyond MAE/MASE, evaluation includes decision-aligned diagnostics motivated by a structural trap: MAE-family metrics are minimized by conditional *medians*, while the downstream optimizer consumes conditional *means* and captaincy rewards the upside tail. The diagnostics - RMSE (mean-aligned), bias, total calibration (ratio of predicted to actual total points), within-gameweek Spearman rank correlation, and top-1 capture (how often the highest-ranked player actually scored most) - make that gap visible instead of letting MASE alone select median-flattening models.

A parallel probabilistic view - per-position LightGBM quantile regressors (p10/p50/p90, crossing repaired by row-wise sorting) - produces a prediction interval per player-gameweek, because two players with equal expected points are not equal decisions once the 2x captain multiplier rewards upside. It does not yet feed the optimizer.

4.3 Evaluation design

Three layers, in increasing decision-relevance:

1. **Static split.** Train on GW ≤ 152, evaluate on GW153-183 - the window the original LSTM was validated on, keeping every comparison on one fixed window.

2. **Walk-forward validation.** For each gameweek from a start point, train strictly on earlier data and predict that gameweek only - many genuinely out-of-sample evaluations rather than one.
3. **Actual-points backtest (gold standard).** Walk-forward predictions run through the MILP over GW153-183; resulting squads scored with realized points. An accuracy gain that does not survive contact with the optimizer is not an improvement (Section 5.4).

Point forecasts are scored with MAE and, primarily, MASE (Hyndman & Koehler, 2006) - MAE relative to the in-sample error of a naive last-gameweek forecast, the scale fit on training data only; MASE < 1 beats the naive floor regardless of the target's scale. Probabilistic forecasts are scored with pinball loss and [p10, p90] coverage against its nominal 0.80.

4.4 Squad optimization

The MILP follows Kristiansen et al. (2018): 15-player squad (2/5/5/3) under budget, max 3 per club, 11-player lineup under formation constraints, captain/vice-captain, limited free transfers with a 4-point penalty per extra, and one-shot chip logic (two wildcards, free hit, bench boost, triple captain), solved as a rolling horizon with only the first gameweek's decision locked in. Venter & van Vuuren (2024) - whose formulation matches almost exactly - attribute their strong case study (top 4.08% of 8.24M managers) to lookahead and forecast quality, not optimizer sophistication; hence this project's focus on forecasting.

5 Results

5.1 Backtest lineage

All systems scored by realized points over 2024-25 GW1-31, identical optimizer. Rows above the line predate the blend-weight leakage fix (weights were fit inside the evaluated window) and are flattered by it; rows below use honest, strictly-prior weight fitting and the full current feature set:

System	**Actual points**
Old LSTM + MILP (thesis)	1526
LightGBM + MILP, 4-season history	1811
Ensemble + MILP, 4-season history	1900
Ensemble + MILP, 6-season history	1966
1. fixture/minutes features	1880[1]
NNLS blend, honest weights, current features	1869
CatBoost-only, honest weights, current features	1856
CatBoost + level recalibration	1800[2]
CatBoost-only, Optuna-tuned hyperparameters	**2107**

Four lessons. Extending history from four to six seasons was an unambiguous win (1,900 $\rightarrow$ 1,966 like-for-like). The honest re-baseline shows the 1,966 headline carried roughly 100 points of weight-leakage-and-feature-drift inflation; 1,870 was the clean untuned baseline. The blend-vs-CatBoost gap (13 points, 0.7%, one window) is noise: the accuracy winner and the calibration

[1] Regression despite improved forecast accuracy (Section 5.4) - the first demonstration that MASE gains need not translate into points.

[2] A theory-driven fix that failed honestly: rescaling each position's deflated forecast level reshuffled cross-position budget allocation and cost 56 points (Section 5.3).

winner build equally good squads, so CatBoost is kept for simplicity and cost, not superiority. And **hyperparameter tuning was worth more than every other intervention combined**: 30 Optuna trials per position (time-ordered CV, tuned strictly on pre-window data) lifted realized points 1,856 → 2,107 (+13.5%) - the tuned parameters uniformly prefer slower learning rates (0.013 vs the hand-set 0.05) and shallower trees, i.e. the defaults were overfitting. **2,107 is the new number to beat**, with the standing caveats: one window, one seed, 30 trials.

5.2 Forecasting-technique experiments

Eight simple and econometric baselines - rolling mean, naive drift, SES, Holt, Theta, Croston, pooled AR(1), per-player ARIMA(1,0,1), and an empirical-Bayes shrinkage of player means toward position means - were each tested per position. None beats the ML models anywhere; SES is the best simple method, Croston and EB-shrinkage the weakest (full tables in `RESEARCH_LOG.md`). Plain OLS, the index, beats every simple baseline; the ML models in turn beat OLS.

5.3 Expanded model registry

Six techniques were added in one round (LinearSVR, RBF SVR, XGBoost, CatBoost, PLS, EB-shrinkage) and evaluated on the standard static split:

Position	CatBoost	LinearSVR	12-member ensemble	previous ensemble
GK	**0.513**	0.513	0.534	0.588
DEF	**0.705**	0.713	0.747	0.799
MID	**0.731**	0.751	0.806	0.799
FWD	**0.849**	0.869	0.853	0.959

MASE, GW153-183 static split. CatBoost uses MAE loss; ensemble scored on a genuine holdout half-window.

Two findings. First, **CatBoost with MAE loss is the best single model at every position**, by a wide margin over everything preceding it - consistent with the hypothesis (raised by LinearSVR's strong showing) that MAE-aligned training objectives suit a zero-inflated target with outlier hauls better than squared error. PLS and XGBoost added nothing (retained as near-zero-weight members). Second, **the blended ensemble now trails standalone CatBoost everywhere** - with 12 collinear members, NNLS overfits its half-window weight fit (the MID blend failed to select CatBoost at all).

That second finding was settled by two independent tests. A 20-gameweek walk-forward head-to-head (GW169-226, members retrained each step, weights fit strictly before the window) had CatBoost-only beat the blend at every position - weighted MASE 0.684 vs 0.761 - and the combination bake-off on the static window agreed, with equal-weight top-k second and ridge stacking last:

Position	CatBoost only	top-k (k=3)	NNLS	ridge
GK	**0.502**	0.517	0.547	0.551
DEF	**0.689**	0.705	0.782	0.801
MID	**0.702**	0.720	0.797	0.780
FWD	**0.788**	0.805	0.884	0.887

Combination bake-off: MASE on identical held-out rows (GW168-183 eval half). The best single member beats every combination scheme - Clemen's (1989) forecast-combination result in the wild.

Production forecasting is therefore CatBoost-per-position, chosen by the bake-off on every training run rather than hard-coded - if a future change makes a combination win again, production follows automatically.

The decision-aligned diagnostics add a critical nuance: **CatBoost is the best ranker but the worst calibrated**. It captures the actual top scorer far better than the blend (top-1 capture 0.57 vs 0.42 at MID; higher at every position) with near-equal rank correlation - yet its bias is -0.32 to -0.60 points per player-gameweek and its forecasts sum to only 44-63% of the points actually scored, exactly the median-flattening the MAE loss predicts on a zero-inflated target. Ranking quality is what transfers and captaincy consume, so this is the right model to keep - but the MILP also makes *absolute-scale* decisions (a 4-point transfer penalty, chip thresholds), which deflated forecasts could in principle distort. The theory-driven fix - a per-position scalar (sum of actual over sum of predicted, fit on the same pre-window holdout as the weights) - was implemented and tested honestly, and *failed*: realized points fell 1,856 → 1,800. The predicted mechanism (suppressed transfers) did not materialise (both runs made 1 transfer per gameweek); instead the unequal scalars (DEF ×2.02 vs GK ×1.39) reshuffled the budget across positions, and that reallocation cost points. The recalibration remains available behind a flag as a documented negative result, off by default (Section 5.1).

A further honest null result: the 37 new features added alongside this batch (opponent form, EWMA form, per-90 rates, xP) left CatBoost's full-window MASE essentially unchanged (0.513→0.517 GK, 0.705→0.706 DEF, 0.731→0.730 MID, 0.849→0.847 FWD). They are retained - the underlying semantic corrections are correctness fixes regardless, and no hyperparameter tuning has yet been attempted that might exploit them - but they have not yet earned their keep on accuracy alone.

5.4 Fixture/minutes features: better forecasts, worse squads

Adding fixture-difficulty and minutes-projection features improved MASE at every position (DEF most, consistent with clean-sheet dependence) - yet realized backtest points *fell* 1,966 → 1,880. Leading hypothesis: the features smooth the mean forecast toward safe, nailed players, and the MILP - blind to variance - loses the high-ceiling captaincy differentials that drive realized hauls. The features are committed but not declared a net win; resolving this (noise-check on a second window, probabilistic-upside captaincy) is a top open item. The probabilistic module's [p10, p90] coverage is 0.88-0.93 against nominal 0.80 - somewhat wide, usable for relative risk ranking; conformal calibration is a possible refinement.

5.5 Exploratory branches

Two isolated branches, neither merged: a Bayesian belief-state MDP manager after Matthews et al. (2012) - 1,083 points myopic / 847 Q-learning where the production system scored 1,900, expected given documented simplifications - and a per-player model-selection plan (feasibility capped by 35% of live players having no prior-season history; recommendation: single-position pilot).

6 Methodological limitations

An internal code-and-concept review (2026-07-04) found four errors; all four were fixed the same day. One further limitation emerged from the diagnostics added afterwards. Recorded here because some published numbers predate the fixes:

1. **Production forecast level is miscalibrated (accepted, correction tested and rejected).** The production CatBoost models under-predict the aggregate point level by 40-55% (a structural consequence of MAE loss on a zero-inflated target) while ranking players better than any alternative. The designed fix - per-position holdout rescaling - was tested and *lost* 56 realized points (Section 5.3, Section 5.1), so the miscalibration is currently accepted as harmless-in-practice for this optimizer. It remains a live concern for any future use of the forecasts where absolute scale matters (e.g. chip-timing heuristics tuned in points).
2. **Blend-weight leakage (fixed, re-baselined).** Backtest predictions formerly reused blend weights fit inside the backtest window itself. Weights are now always fit on a window strictly before whatever is predicted. The re-baseline quantified the damage: the honest figure is 1,869 against the leaky 1,966 (Section 5.1), so roughly 100 points of the old headline were leakage-and-feature-drift, not skill.
3. **Live-mode staleness (fixed).** The weekly driver formerly reused each player's last played row, so live fixture-difficulty described last week's fixture and rolling form excluded the most recent match; API-vs-dataset team-name mismatches also silently dropped some teams' players. Live predictions now use per-gameweek API fixture difficulties and a synthetic next-gameweek row whose features include everything played. Backtests were never affected.
4. **Index-check asymmetry (fixed).** The ensemble-vs-OLS verdict now compares both on the same held-out rows.
5. **Optimizer rule currency (open).** The MILP caps banked free transfers at 2 (now 5 in FPL) and assumes sales at market value (real FPL: purchase price plus half profit). Fine for historical comparison; wrong for live 2026-27 play. Deliberately deferred - optimizer work is deprioritized in favor of forecasting.

7 References

Croston, J. D. (1972). Forecasting and stock control for intermittent demands. *Journal of the Operational Research Society*, 23(3), 289-303.

Hyndman, R. J., & Koehler, A. B. (2006). Another look at measures of forecast accuracy. *International Journal of Forecasting*, 22(4), 679-688.

Kristiansen, B. K., Gupta, A., & Eilertsen, W. (2018). *Developing a forecast-based optimisation model for Fantasy Premier League* (MSc thesis). Norwegian University of Science and Technology, Trondheim.

Matthews, T., Ramchurn, S. D., & Chalkiadakis, G. (2012). Competing with humans at fantasy football: Team formation in large partially-observable domains. *Proceedings of the Twenty-Sixth AAAI Conference on Artificial Intelligence*, 1394-1400.

Venter, V., & van Vuuren, J. H. (2024). An optimisation approach towards soccer Fantasy Premiere League team selection. *ORiON*, 40(1), 69-107.